

Concurrent Priority Queues

Seminar in Algorithms, 2013W

Jakob Gruber, 0203440

September 23, 2015

Introduction I

Priority Queues (PQs):

- ▶ Standard abstract data structure
- ▶ Used widely in algorithmics, operating systems, task scheduling, etc
- ▶ Interface consists of two $O(\log n)$ operations:

```
void Insert(pq_t *pq, key_t k, value_t v)
bool DeleteMin(pq_t *pq, value_t *v)
```

- ▶ Typical backing data structures: heaps & search trees

Introduction II

- ▶ In the past decade, processor clock speeds have remained the same, trend towards multiple cores
- ▶ New data structures required to take advantage of concurrent execution
- ▶ The topic of this presentation: efficient concurrent PQs
- ▶ Fine-grained locking → Lock-free → Relaxed data structures

Concepts and Definitions

Safety conditions: nothing bad has happened yet

- ▶ *Linearizability*: operations appear to take effect at a single point in time, the linearization point
- ▶ *Quiescent consistency*: in a period of quiescence, semantics equivalent to some sequential ordering

Concepts and Definitions

Liveness conditions: something good eventually happens

- ▶ *Lock-freedom*: at least a single process makes progress at all times
- ▶ *Wait-freedom*: every process finishes in a bounded number of steps

Concepts and Definitions

Miscellaneous

- ▶ *Disjoint-access parallelism*: how well a data structure handles concurrent use by multiple threads within disjoint areas
- ▶ Synchronization primitives:
 - ▶ Compare-And-Swap (CAS), Fetch-And-Add (FAA), Fetch-And-Or (FAO), Test-And-Set (TAS)
 - ▶ Double-Compare-And-Swap (DCAS), Double-Compare-Single-Swap (DCSS)

```
bool CAS(T *ptr, T *expected, T value) {  
    if (*ptr == *expected) {  
        *ptr = value;  
        return true;  
    } else {  
        *expected = *ptr;  
        return false;  
    }  
}
```

Related Work

- ▶ Non-standard synchronization primitives
 - ▶ Liu and Spear: Array-based PQ with ExtractMany
 - ▶ Israeli and Rappoport: Wait-free PQ
- ▶ Bounded range priorities
 - ▶ Shavit and Zemach: Combining funnels & bins
- ▶ Strict PQs
 - ▶ Hunt et al.: Fine-grained locking heap
 - ▶ Shavit and Lotan: First SkipList-based PQ
 - ▶ Sundell and Tsigas: First lock-free PQ
 - ▶ Lindén and Jonsson: Minimizes contention
- ▶ Relaxed data structures
 - ▶ Kirsch, Lippautz, and Payer: k-FIFO queues
 - ▶ Wimmer et al.: k-PQ
 - ▶ Alistarh et al.: SprayList

Fine-grained Locking Heaps

Hunt et al.

- ▶ Naive PQ parallelization: single global lock \rightarrow sequential bottleneck
- ▶ A first improvement: fine-grained locking using a lock per node
- ▶ Galen C Hunt et al. “An efficient algorithm for concurrent priority queue heaps”. In: *Information Processing Letters* 60.3 (1996), pp. 151–157

Fine-grained Locking Heaps

Hunt et al.: Innovations

- ▶ One lock per node, *but* additionally a global lock protecting the heap's size variable
- ▶ Insertions bottom-up, deletions top-down to reduce contention
- ▶ Successive insertions take disjoint paths towards the root

Fine-grained Locking Heaps

Hunt et al.: Limitations

- ▶ A global lock remains
- ▶ Heap is statically allocated
- ▶ Frequent complex heap reorganization
- ▶ Disjoint-access breaks down at high traffic levels
- ▶ Inherent PQ bottleneck at the minimal node
- ▶ Benchmarks show only limited scalability up to a low thread count

Lock-free Priority Queues

SkipLists

- ▶ Modern concurrent PQs are mostly based on Pugh's SkipList (SL)
- ▶ Probabilistic ordered search structure, insertions and deletions in expected $O(\log n)$ time
- ▶ No reorganizations
- ▶ Simple implementation
- ▶ Excellent disjoint-access properties

Lock-free Priority Queues

SkipLists

- ▶ Collection of linked lists with corresponding levels
- ▶ Lowest list contains all items, higher lists are shortcuts
- ▶ Insert chooses a level according to geometric distribution

```
struct slist_t {  
    size_t max_level;  
    node_t head[max_level];  
};  
  
struct node_t {  
    key_t key;  
    value_t value;  
    size_t level;  
    node_t *next[level];  
};
```

Lock-free Priority Queues

Shavit and Lotan

- ▶ First SkipList-based PQ
- ▶ Nir Shavit and Itay Lotan. “Skiplist-based concurrent priority queues”. In: *Proceedings of the 14th International Symposium on Parallel and Distributed Processing (IPDPS)*. IEEE. 2000, pp. 263–268
- ▶ Initially lock-based (linearizable), lock-free (quiescently consistent) variant published in 2008
- ▶ Maurice Herlihy and Nir Shavit. *The Art of Multiprocessor Programming, Revised Reprint*. Elsevier, 2012

Lock-free Priority Queues

Shavit and Lotan




Figure: The Shavit and Lotan PQ (Image source: [9])

Lock-free Priority Queues

Shavit and Lotan

- ▶ Items are considered in the list once inserted on bottom level
- ▶ Again, insertions bottom-up and deletions top-down
- ▶ Deletions are split
 - ▶ Logical deletion sets a deleted flag
 - ▶ Physical deletion performs actual pointer manipulations
- ▶ DeleteMin attempts to logically delete the head node. On success: delete physically & return node. Otherwise, continue with next node.
- ▶ Insert is equivalent to the SL insertion

Lock-free Priority Queues

Shavit and Lotan

- ▶ Quiescently consistent, but not linearizable
 - ▶ Slow thread A suspended at deleted key k while in `DeleteMin`
 - ▶ Fast thread B first inserts $k - 1$, then $k + 1$
 - ▶ A wakes up and returns $k + 1$
 - ▶ Linearizability would require returning $k - 1$
- ▶ Timestamp mechanism: stamp nodes on successful insertion, `DeleteMin` ignores all stamps earlier than its own starting point
- ▶ Improved scalability, but heavy contention at list head

Lock-free Priority Queues

Sundell and Tsigas

- ▶ First lock-free PQ, linearizable, SL-based, distinct priorities
- ▶ Deletion flag packed into least significant bits of `next` pointers prevent insertion *after* deleted nodes
- ▶ Helping mechanism ensures only a single logically deleted node exists at any time
- ▶ Performs significantly better than the Hunt et al. heap, slightly better than a SkipList protected by a global lock
- ▶ Håkan Sundell and Philippos Tsigas. “Fast and lock-free concurrent priority queues for multi-thread systems”. In: *Proceedings of the 17th International Symposium on Parallel and Distributed Processing (IPDPS)*. IEEE. 2003, 11–pp

Lock-free Priority Queues

Lindén and Jonsson

- ▶ Most efficient strict PQ, SL-based & linearizable
- ▶ Concurrent strict PQ performance limited by contention and CAS failures in DeleteMin → Minimize CAS calls
- ▶ Deleted nodes form prefix of SL, deletion flag packed into next pointer of *previous* node prevents insertion *before* deleted node
- ▶ Most DeleteMin perform logical deletion (1 CAS) only, physical deletion only once a bound is reached
- ▶ Improves upon best previous PQs by up to factor 2
- ▶ Jonatan Lindén and Bengt Jonsson. “A Skiplist-Based Concurrent Priority Queue with Minimal Memory Contention”. In: *Principles of Distributed Systems*. Springer, 2013, pp. 206–220

Lock-free Priority Queues

Lindén and Jonsson




Figure: The Lindén and Jonsson PQ. Concurrent `Insert(3)` and `DeleteMin` operations. Top right: `DeleteMin` succeeds first, `Insert(3)` CAS fails. Bottom right: `Insert(3)` succeeds first, `DeleteMin` returns 3. (Image source: [6])

Relaxed Priority Queues

- ▶ Strict PQs have inherent bottleneck at minimal element
- ▶ To improve further: < 1 CAS per DeleteMin
- ▶ Another approach is to relax semantics, i.e. instead of returning *the* minimal element, return one of the k minimal elements

Relaxed Priority Queues

Alistarh et al.

- ▶ Relaxed SL-based PQ, safety properties unclear
- ▶ DeleteMin returns one of the $O(P \log^3 P)$ elements
- ▶ Degrades to random-remove if the PQ is small compared to thread count P (for $P = 80$, $P \log^3 P \approx 7000$)
- ▶ DeleteMin performs random walk: starting at the list head on some level l , repeatedly follow a randomized number of `next[1]` pointers, descend a randomized number of levels
- ▶ Parameters are chosen s.t. each element within the walk has approximately equal probability of being returned
- ▶ Benchmarks show scalability comparable to a random-remove Delete up to at least 80 threads
- ▶ Dan Alistarh et al. *The SprayList: A Scalable Relaxed Priority Queue*. Tech. rep. 2014

Relaxed Priority Queues

Wimmer et al.

- ▶ First relaxed linearizable PQs: we discuss the hybrid k -PQ which, provides a bound of kP missed elements in `DeleteMin`
- ▶ Consists of a list of globally visible elements, and per thread: a local item list, and a local PQ
- ▶ Insert accesses only the local structures as long as guarantees are not violated, otherwise the global list is updated the the local view is synchronized
- ▶ `DeleteMin` pops the local queue if it is non-empty; otherwise spy, i.e. copy elements from another thread's local list
- ▶ Very good scalability up to 10 threads, further limited gains with rising thread counts
- ▶ Martin Wimmer et al. "Data Structures for Task-based Priority Scheduling". In: *arXiv preprint arXiv:1312.2501* (2013)

Lock-free Priority Queues

Wimmer et al.

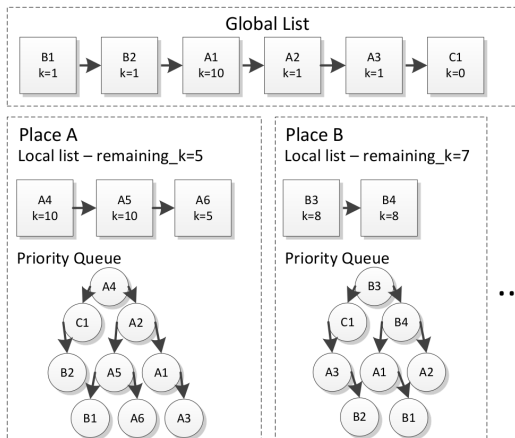


Figure: The Wimmer et al. hybrid k-PQ. (Image source: [12])

Results

- ▶ Benchmarking results from selected papers & own results
- ▶ We present throughput, i.e. the number of operations performed per second
- ▶ Each thread repeatedly chooses uniformly at random between Insert and DeleteMin

Results

Lindén and Jonsson

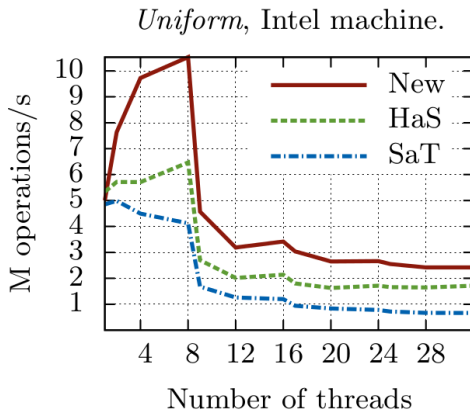


Figure: GCC 4.7.2, 32-core Intel Xeon E5-4650 @ 2.7 GHz. Initial PQ size unknown. *New*: Lindén and Jonsson, *HaS*: Shavit and Lotan, *SaT*: Sundell and Tsigas. (Image source: [6])

Results

Gruber




Figure: GCC 4.8.2, 80-core Intel Xeon E7-8850 @ 2.0 GHz. PQ initialized with 2^{15} elements.

Results

Alistarh et al.

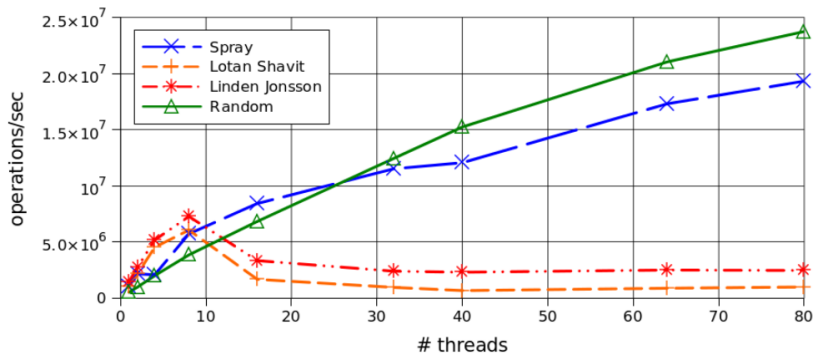


Figure: GCC version unknown, 80-core Intel Xeon E7-4870 @ 2.4 GHz.
PQ initialized with 10^6 elements. (Image source: [1])

Conclusion

- ▶ Parallelizing PQs is hard
- ▶ SkipLists currently dominate practical implementations
- ▶ Main limiting factor are list head accesses in DeleteMin
- ▶ Lindén and Jonsson PQ is state of the art in strict semantics
- ▶ Much potential remains in relaxed PQs and data structures in general → future research

Conclusion

Questions?

References I

Dan Alistarh et al. *The SprayList: A Scalable Relaxed Priority Queue*. Tech. rep. 2014.

Maurice Herlihy and Nir Shavit. *The Art of Multiprocessor Programming, Revised Reprint*. Elsevier, 2012.

Galen C Hunt et al. “An efficient algorithm for concurrent priority queue heaps”. In: *Information Processing Letters* 60.3 (1996), pp. 151–157.

Amos Israeli and Lihu Rappoport. “Efficient wait-free implementation of a concurrent priority queue”. In: *Distributed Algorithms*. Springer, 1993, pp. 1–17.

Christoph M Kirsch, Michael Lippautz, and Hannes Payer. *Fast and scalable k-fifo queues*. Tech. rep. 2012-04, Department of Computer Sciences, University of Salzburg, 2012.

References II

Jonatan Lindén and Bengt Jonsson. “A Skiplist-Based Concurrent Priority Queue with Minimal Memory Contention”. In: *Principles of Distributed Systems*. Springer, 2013, pp. 206–220.

Yujie Liu and Michael Spear. “A lock-free, array-based priority queue”. In: *ACM SIGPLAN Notices* 47.8 (2012), pp. 323–324.

William Pugh. “Skip lists: a probabilistic alternative to balanced trees”. In: *Communications of the ACM* 33.6 (1990), pp. 668–676.

Nir Shavit and Itay Lotan. “Skiplist-based concurrent priority queues”. In: *Proceedings of the 14th International Symposium on Parallel and Distributed Processing (IPDPS)*. IEEE, 2000, pp. 263–268.

References III

Nir Shavit and Asaph Zemach. “Scalable concurrent priority queue algorithms”. In: *Proceedings of the 18th ACM Symposium on Principles of Distributed Computing (PODC)*. ACM. 1999, pp. 113–122.

Håkan Sundell and Philippas Tsigas. “Fast and lock-free concurrent priority queues for multi-thread systems”. In: *Proceedings of the 17th International Symposium on Parallel and Distributed Processing (IPDPS)*. IEEE. 2003, 11–pp.

Martin Wimmer et al. “Data Structures for Task-based Priority Scheduling”. In: *arXiv preprint arXiv:1312.2501* (2013).